

HDS06 - Designing and Querying a Clinical Trial Participant Registry

Westbrook University Hospitals NHS Trust, Clinical Trial Participant Registry

Malik J. Lainsbury | [github](#)

2026-07-11



UNIVERSITY OF
CAMBRIDGE

Professional and
Continuing Education

Table of contents

Introduction	2
Part 1: Database Design	3
Entity Relationship Diagram	3
Table Design Choices	4
Column Design Choices	5
Relationships Between Tables	7
Part 2: SQL	9
INSERT statements	9
SQL Queries	14
Part 3: Reflection	20
Why no participant names?	20
A situation where the data could mislead	20
Why consent records are append-only, and how to enforce it	21
References	22
Appendix A: Data dictionary	23
ClinicalTrial	23
HospitalSite	24
Participant	24
InactivationReason	24
TrialSite (link table)	25
ConsentVersion	25
Enrolment	25
ConsentSigning (append-only)	26

Introduction

This report designs a fictional relational database (DB) for the Westbrook University Hospitals NHS Trust Clinical Trial Participant Registry, implements it in Microsoft SQL Server (T-SQL), populates it with fictional data and answers seven clinical questions via SQL queries.

Note

Reproducing the results. Please run `script/HDS_DB_LAINSBURY_2026.sql`. This is the requested single self-contained script that creates the database, tables, data, trigger and queries in order. It ran without errors from start to finish on my machine. Each query returns one result grid, screenshotted below. I used VS Code with the mssql extension on a local SQL Server.

Part 1: Database Design

Entity Relationship Diagram

My database schema (dbo.) has eight tables: the five reference tables provided as .csv files, plus three of my own. The table **TrialSite** resolves the trial-site many-to-many, **Enrolment** records a participant's participation in a trial at a site and **ConsentSigning** is the append-only audit log of consent events. The ERD is shown in crow's-foot notation in Figure 1.




Figure 1: Entity Relationship Diagram (crow's-foot notation).

Table Design Choices

Foreign and Unique Key Design Choices

Every table uses a surrogate INT IDENTITY primary key for internal joins, plus one or more **UNIQUE keys (or business keys)** for the values clinical staff would actually recognise.

Every foreign key targets the surrogate key, never the business key. This renders the inserts slightly more verbose but is a personal choice for data integrity and maintainability. This study uses externally issued codes such as the ISRCTN number. If they were to change or input wrongfully, they can then be corrected in one parent row without cascading. The unique keys are mainly used in INSERT look-up subqueries. In the real world, the business keys would be known to the data-entry staff and the look-ups would be hidden behind a user interface/form or via a staging table.

Table Descriptions

Table 1: Table descriptions and key design choices	
Table	Purpose and key design choices
ClinicalTrial	One row per trial. PK is <code>trial_id</code> (IDENTITY). UNIQUE <code>registration_number</code> (ISRCTN number) is the externally recognised natural key. Children tables reference <code>trial_id</code> , never the ISRCTN number.
HospitalSite	One row per hospital site. the PK is <code>site_id</code> . the UNIQUE <code>site_code</code> (e.g. WBK-GEN) are the operational natural key. Children tables reference <code>site_id</code> .
Participant	One unique pseudonymised participant per row, no names held . the PK is <code>participant_id</code> (IDENTITY).
InactivationReason	Controlled vocabulary of reasons a participant leaves a trial, so exits report consistently across trials. PK <code>reason_id</code> (IDENTITY). The UNIQUE <code>reason_code</code> is the stable report token.
TrialSite	Link table resolving the trial-site many-to-many. It also records when each site opened/closed for recruitment on that trial. It uses a composite PK (<code>trial_id</code> , <code>site_id</code>). The pairing is the row's identity, so no surrogate is used (unlike every other table). FKs to ClinicalTrial and HospitalSite. Both FKs reference the parents' surrogate keys.
ConsentVersion	Every approved version of a trial's consent form, including wording and effective date. Any old consent versions is never overwritten (under the principle of always add rather than overwrite). the PK is <code>consent_version_id</code> . there can only be one of the same version number per trial:UNIQUE (<code>trial_id</code> , <code>version_number</code>). The FK references <code>trial_id</code> , not the ISRCTN number to follow the same design pattern as other tables.

Table	Purpose and key design choices
Enrolment	A participant enrolled in a trial at a specific site. It carries (<code>trial_id</code> , <code>site_id</code>) as a composite FK straight to <code>TrialSite</code> , so trial and site are captured together and a participant can only enrol where the trial actually runs. The alternative would of been a surrogate <code>trial_site_id</code> pointer that would enforce the same rule but force every query to hop through <code>TrialSite</code> . By storing the pair keeps enrolment rows self-describing without denormalising (the composite FK keeps the pair consistent). Exit is recorded without deleting the row. PK is <code>enrolment_id</code> . The UNIQUE (<code>participant_id</code> , <code>trial_id</code> , <code>site_id</code> , <code>enrolment_date</code>) is a duplicate-entry guard to follow the brief: the same enrolment cannot be keyed twice on the same day (the likeliest failure when replacing spreadsheet data entry), while a participant who left a trial can still legitimately re-enrol later. Concurrent enrolment is prevented separately by the filtered index.
ConsentSigning	Append-only log of every consent signing event. It holds when, which version, who witnessed. The PK is <code>consent_signing_id</code> . FKs are to <code>Enrolment</code> and <code>ConsentVersion</code> . UNIQUE (<code>enrolment_id</code> , <code>consent_version_id</code> , <code>signed_date</code>) treats a same-day duplicate row as a keying error while keeping a genuine re-signing of the same version on a later date (a damaged paper form re-executed) recordable. The brief requires every signing event to be recorded. The table is protected by an append-only trigger. Both FKs reference surrogate keys. the CHECK requires <code>witnessed_by</code> to contain at least a forename and surname for tracking purposes.

Column Design Choices

Column constraints

- Phase/status/sex use `CHECK ... IN (...)` to force controlled vocabularies and clean categorical data for reporting. **Caveat:** `sex_at_birth` is interpreted as biological sex at birth (not gender identity) and captures only two values to keep the schema simple. In a production DB this would need to be reviewed against best practice and local healthcare inclusivity guidance.
- I use date checks (`end_date >= start_date`, `closed_date >= opened_date`, `inactivation_date >= enrolment_date`, `date_of_birth < registration_date`) throughout to prevent logically impossible records. The first three use `>=` because a same-day pair is legitimate. The birth-date check is deliberately strict (`<`) as a participant cannot register with the research office on or before the day they were born.
- A `CHECK (version_number > 0)` on `ConsentVersion` rejects zero or negative version numbers, keeping the version sequence meaningful.
- The paired-null check on `Enrolment` guarantees a leaving date never appears without a

reason, and vice-versa.

- A `CHECK` on `witnessed_by` requires at least two words (forename + surname), so a single word or an initial alone cannot be recorded as a witness. The characters themselves are deliberately unrestricted so that names such as José or O'Brien are stored exactly as written. This matters for an auditable consent record.
- I used `NVARCHAR(MAX)` for consent wording because it can be long and multi-paragraph. The other text fields are sized to fit the expected content.
- A filtered unique index (`UX_Enrolment_OneActivePerParticipant`) on `Enrolment.participant_id` WHERE `inactivation_date` IS NULL enforces the business rule that a participant may be actively enrolled in only one trial at a time as requested in the brief. This allows unlimited historical enrolments while preventing concurrent participation.
- I added a trigger to enforce the append-only rule on `ConsentSigning`. It rejects any `UPDATE` or `DELETE` operation with a clear error message (51000), while leaving `INSERT` free. This guarantees the rule regardless of which application or user connects, which a permission grant alone would not.

NULL and NOT NULL Decisions

The two tables below cover most column in the schema. `IDENTITY` primary keys and `TrialSite`'s composite-PK columns are NOT NULL by definition and are omitted from below tables.

Table 2: Columns allowing NULL values

Table	Column	Rationale
ClinicalTrial	<code>end_date</code>	If NULL, the trial is still running. Current trials have no end date <i>yet</i> .
TrialSite	<code>closed_date</code>	The site can still be open for recruitment on that trial.
Enrolment	<code>inactivation_date</code> , <code>inactivation_reason_id</code>	This is a “meaningful absence”: the participant is still active in the trial. The pair must be NULL (or populated) together therefore the paired-null <code>CHECK</code> keeps them in step.

Table 3: Columns strictly NOT NULL

Table	Column	Rationale
ClinicalTrial	<code>start_date</code> , <code>status</code> , <code>phase</code>	A trial cannot exist without these columns populated.
TrialSite;	<code>opened_date</code> ;	Every record's timeline must begin somewhere.
Enrolment;	<code>enrolment_date</code> ;	Unlike their end-date above, an absent value
ConsentVersion	<code>effective_from</code>	would carry no meaning.

Table	Column	Rationale
Participant	date_of_birth, registration_date, sex_at_birth	Required for eligibility, age and demographic reporting.
ConsentSigning	signed_date, witnessed_by	A consent event without a date or witness is not a valid consent event.
ConsentVersion	wording_text	A consent version is meaningless without its wording.
Enrolment; ConsentSigning; ConsentVersion	FK columns: participant_id, trial_id, site_id; enrolment_id, consent_version_id; trial_id	A child row cannot exist without its parents. Only exception is with inactivation_reason_id, the one deliberately nullable FK, paired with inactivation_date.
ClinicalTrial; HospitalSite; Participant; InactivationReason; ConsentVersion	Unique keys: registration_number; site_code; study_code; reason_code; version_number	These UNIQUE keys are the look-up targets of the INSERT subqueries. A NULL here would break that pattern and the insertion would fail.
ClinicalTrial; HospitalSite; InactivationReason	Descriptive fields: trial_name; site_name, city, country; reason_description	A reference row without its human-readable label is unusable in reports, it is a reporting requirement.

Relationships Between Tables

Table 4: Relationships between tables

Tables	Type / FK	What it represents
ClinicalTrial – HospitalSite	M:N via TrialSite (FK trial_id, site_id)	A trial runs at many sites and a site hosts many trials. The link table also carries the open/close recruitment dates.
ClinicalTrial -> ConsentVersion	1:M (FK trial_id)	One trial accumulates many consent-form versions over its life as the protocol changes.
Participant -> Enrolment	1:M (FK participant_id)	A participant may have several enrolments over time, only one active at once as defined per brief (filtered index).

Tables	Type / FK	What it represents
TrialSite -> Enrolment	1:M (composite FK <code>trial_id, site_id</code>)	Each enrolment happens at exactly one trial-site pairing that the TrialSite link table confirms actually exists.
InactivationReason -> Enrolment	1:M optional (FK <code>inactivation_reason_id</code>)	A standardised reason classifies many exits; active enrolments reference none.
Enrolment -> ConsentSigning	1:M (FK <code>enrolment_id</code>)	One enrolment can have several signing events (original + re-consent after amendment).
ConsentVersion -> ConsentSigning	1:M (FK <code>consent_version_id</code>)	Each signing records which version was signed; one version is signed by many participants.

Many-to-many resolution. The only many-to-many relationship (ClinicalTrial - HospitalSite) is resolved with the **TrialSite** junction table, using the composite primary key (`trial_id, site_id`). The pairing itself is the row's identity.

Part 2: SQL

INSERT statements

Every INSERT that references another table (e.g. Enrolment needs a trial_id) looks the ID up by the human-recognisable unique key. There is no use of hard-coded IDENTITY integers anywhere. It made the script more verbose but is best practice for maintainability, as discussed in the key design choices above. The added data is fictional and does not represent any real trial, site or participant.

Note

As requested in the brief, only the first INSERT per table is shown here; the rest are in `script/HDS_DB_LAINSBURY_2026.sql`.

ClinicalTrial

The three trials come straight from the reference data provided (`data/ClinicalTrial.csv`). No FK is needed so the values are inserted directly. SQL Server generates the IDENTITY `trial_id`.

```
INSERT INTO dbo.ClinicalTrial (registration_number, trial_name, phase, status,
    ↪ start_date, end_date)
VALUES ('ISRCTN10234567', N'CARDIOPROTECT: Cardiac Rehabilitation in Post-MI
    ↪ Patients',
        'III', 'active', '2023-03-01', NULL);
```

HospitalSite

The three hospital sites are from the reference data (`data/HospitalSite.csv`), inserted directly with no look-up needed.

```
INSERT INTO dbo.HospitalSite (site_code, site_name, city, country)
VALUES ('WBK-GEN', N'Westbrook General Hospital', N'London', N'England');
```

Participant

All five pseudonymised participants are from the reference data (`data/Participant.csv`). Only the study code identifies them, no names are held.

```
INSERT INTO dbo.Participant (study_code, date_of_birth, sex_at_birth,
    ↪ registration_date)
```

```
VALUES ('WBK-001', '1958-04-12', 'Male', '2021-11-03');
```

InactivationReason

The six standardised reasons are from the reference data (data/InactivationReason.csv). They form the controlled vocabulary that **Enrolment** references when a participant leaves a trial.

```
INSERT INTO dbo.InactivationReason (reason_code, reason_description)
VALUES ('consent_withdrawn', N'Participant withdrew their consent to take part
↳ in the trial');
```

TrialSite

This data is entirely my own, the reference data contains no information about which trial runs at which site. I linked CARDIOPROTECT and BREATHE to two sites each and MOBILISE to one. Both FKs are looked up by unique key (registration number and site code).

```
INSERT INTO dbo.TrialSite (trial_id, site_id, opened_date, closed_date)
VALUES (
  (SELECT trial_id FROM dbo.ClinicalTrial WHERE
↳ registration_number='ISRCTN10234567'),
  (SELECT site_id FROM dbo.HospitalSite WHERE site_code='WBK-GEN'),
  '2023-03-01', NULL);
```

ConsentVersion

The four consent versions are from the reference data (data/ConsentVersion.csv), which identifies each trial by its registration number. I resolve that to **trial_id** with a subquery rather than hardcoding the IDENTITY value.

```
INSERT INTO dbo.ConsentVersion (trial_id, version_number, effective_from,
↳ wording_text)
VALUES (
  (SELECT trial_id FROM dbo.ClinicalTrial WHERE
↳ registration_number='ISRCTN10234567'),
  1, '2023-03-01',
  N'I agree to take part in the CARDIOPROTECT study. ... I may withdraw at any
↳ time without affecting my medical care.');
```

Enrolment

This data is entirely my own. I created eight enrolments covering active, completed, withdrawn and deceased participants. They are spread across two sites (leaving one site with 0 enrolments). All three FKs are resolved by subquery from the study code, registration number and site code.

```
INSERT INTO dbo.Enrolment (participant_id, trial_id, site_id, enrolment_date,
↳ inactivation_date, inactivation_reason_id)
VALUES (
  (SELECT participant_id FROM dbo.Participant WHERE study_code='WBK-001'),
  (SELECT trial_id FROM dbo.ClinicalTrial WHERE
↳ registration_number='ISRCTN10234567'),
```

```
(SELECT site_id FROM dbo.HospitalSite WHERE site_code='WBK-GEN'),  
'2023-03-15', NULL, NULL);
```

ConsentSigning

Data is entirely my own. I added nine signing events, one per enrolment plus WBK-001's re-consent to CARDIOPROTECT v2 after the protocol amendment. The enrolment is found via participant and trial, the version via trial and version number.

```
INSERT INTO dbo.ConsentSigning (enrolment_id, consent_version_id, signed_date,
    witnessed_by)
VALUES (
    (SELECT e.enrolment_id FROM dbo.Enrolment e
        WHERE e.participant_id = (SELECT participant_id FROM dbo.Participant
            WHERE study_code = 'WBK-001')
        AND e.trial_id = (SELECT trial_id FROM dbo.ClinicalTrial
            WHERE registration_number = 'ISRCTN10234567')),
    (SELECT cv.consent_version_id FROM dbo.ConsentVersion cv
        WHERE cv.trial_id = (SELECT trial_id FROM dbo.ClinicalTrial
            WHERE registration_number = 'ISRCTN10234567')
        AND cv.version_number = 1),
    '2023-03-15', N'Amara Okafor');
```

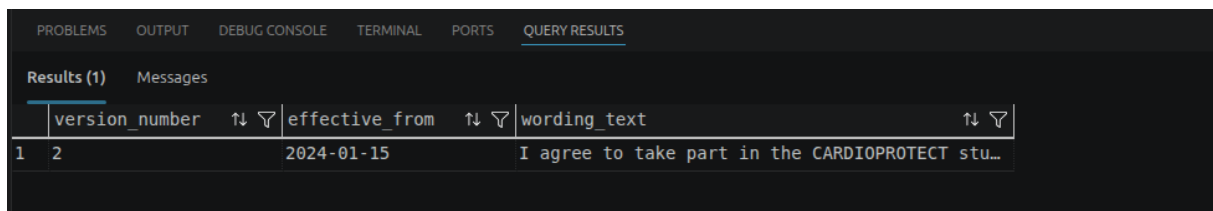
SQL Queries

Query 1: Current consent wording for CARDIOPROTECT

Clinical question. Before conducting a consent discussion with a participant, a research nurse needs to retrieve the current wording of the consent form for the CARDIOPROTECT trial, the version that came into effect most recently. Show the version number, the date it came into effect, and the full wording text.

```
SELECT TOP (1) cv.version_number, cv.effective_from, cv.wording_text
FROM    dbo.ConsentVersion cv
JOIN    dbo.ClinicalTrial ct ON ct.trial_id = cv.trial_id
WHERE   ct.trial_name LIKE 'CARDIOPROTECT%'
ORDER  BY cv.effective_from DESC;
```

Screenshot.



Results (1)			
Messages			
	version_number	effective_from	wording_text
1	2	2024-01-15	I agree to take part in the CARDIOPROTECT stu...

Figure 2: Query 1 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. The query filters to CARDIOPROTECT, orders versions by `effective_from` descending and keeps the top row, so it always returns the latest version even after future amendments.

Data quality: it relies on `effective_from` being accurate. The NOT NULL constraint prevents a missing date, but a wrong-yet-plausible date cannot be caught by any constraint and could show the nurse out-of-date wording.


Query 2: Monthly trial status report

Clinical question. A trial coordinator is preparing a monthly status report. She needs a summary showing each trial, its current status, how long it has been running in days, and whether it has an end date recorded or is still ongoing. Order the results with active trials first, then recruiting, then all others.

```
SELECT ct.trial_name, ct.status,
       DATEDIFF(DAY, ct.start_date, COALESCE(ct.end_date, CAST(GETDATE() AS
    DATE))) AS days_running,
       CASE WHEN ct.end_date IS NULL THEN 'Ongoing' ELSE 'Ended' END AS
    end_date_status
FROM    dbo.ClinicalTrial ct
ORDER  BY CASE ct.status WHEN 'active' THEN 1 WHEN 'recruiting' THEN 2 ELSE 3
    END,
```

```
ct.trial_name;
```

Screenshot.



	trial_name	status	days_running	end_date_status
1	CARDIOPROTECT: Cardiac Rehabilitation in Post-MI P...	active	1227	Ongoing
2	BREATHE: Pulmonary Rehabilitation in COPD	recruiting	769	Ongoing
3	MOBILISE: Early Mobilisation Following Hip Replace...	completed	1460	Ended

Figure 3: Query 2 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. I used COALESCE to substitute today's date for trials with no end date, so "days running" is meaningful for ongoing trials as well as completed ones. I used the CASE expression in ORDER BY to give the custom status priority.


Data quality: "days running" for ongoing trials changes every day the report is run, so it is only a dated snapshot. It would be a great query in a dynamic report such as Power BI or Tableau, but in a static report it is only accurate on the day it is run.

Query 3: Trial history for participant WBK-003

Clinical question. A research nurse is about to meet with participant WBK-003 and wants to see their full trial history: every trial they have ever been enrolled in, when they enrolled, whether they are still active, and if not why they left. Show the trial name, enrolment date, inactivation date, and inactivation reason code, ordered by enrolment date.

```
SELECT ct.trial_name, e.enrolment_date, e.inactivation_date, ir.reason_code
FROM   dbo.Enrolment      e
JOIN   dbo.Participant    p ON p.participant_id = e.participant_id
JOIN   dbo.ClinicalTrial  ct ON ct.trial_id      = e.trial_id
LEFT   JOIN dbo.InactivationReason ir ON ir.reason_id =
      e.inactivation_reason_id
WHERE  p.study_code = 'WBK-003'
ORDER BY e.enrolment_date;
```

Screenshot.



	trial_name	enrolment_date	inactivation_date	reason_code
1	BREATHE: Pulmonary Rehabilitation in COPD	2024-06-10	2024-09-15	consent_withdrawn
2	CARDIOPROTECT: Cardiac Rehabilitation in P...	2024-10-01	NULL	NULL

Figure 4: Query 3 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. The LEFT JOIN to InactivationReason is used because an active enrolment always has a NULL inactivation reason and an INNER join would silently drop it. Additionally, in my design, Enrolment stores trial_id directly (as part of its composite FK to TrialSite). Therefore ClinicalTrial is joined straight off it with no TrialSite hop needed.

Data quality: a participant who left but whose inactivation_reason_id was never populated would show a leaving date with a blank reason. This is why in the creation of the table, I included the paired-null check constraint to prevent this.

Query 4: Audit of currently active enrolments

Clinical question. A clinical research associate is auditing active enrolments across all trials. She needs a report showing each currently **active** participant, alongside the trial they are enrolled in, the hospital site where they are based, and the date they enrolled. Order by trial name then enrolment date.

```
SELECT p.study_code, ct.trial_name, hs.site_name, e.enrolment_date
FROM   dbo.Enrolment     e
JOIN   dbo.Participant    p ON p.participant_id = e.participant_id
JOIN   dbo.ClinicalTrial  ct ON ct.trial_id      = e.trial_id
JOIN   dbo.HospitalSite  hs ON hs.site_id       = e.site_id
WHERE  e.inactivation_date IS NULL
ORDER BY ct.trial_name, e.enrolment_date;
```

Screenshot.



	study_code	trial_name	site_name	enrolment_date
1	WBK-001	CARDIOPROTECT: Cardiac Rehabilitation in P...	Westbrook General Hospital	2023-03-15
2	WBK-003	CARDIOPROTECT: Cardiac Rehabilitation in P...	Westbrook General Hospital	2024-10-01
3	WBK-002	CARDIOPROTECT: Cardiac Rehabilitation in P...	St Katherine's Hospital	2025-01-10
4	WBK-005	CARDIOPROTECT: Cardiac Rehabilitation in P...	St Katherine's Hospital	2025-04-01

Figure 5: Query 4 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. “Active” is derived from inactivation_date IS NULL. Again, as I used a composite FK in my table design, Enrolment stores both trial_id and site_id directly, so ClinicalTrial and HospitalSite are joined straight off it. This avoids the TrialSite hop needed to reconstruct where each active participant sits.


Data quality: the definition of “active” depends on inactivation being recorded promptly as it is derived from a NULL value in the inactivation_date column. A participant who has left but whose exit was not entered would wrongly appear here.

Query 5: Enrolment count per site (2 queries depending on interpretation of the request)

Clinical question. The research director wants to see how many participants have ever been enrolled at each hospital site, including sites that currently have no enrolments recorded. Show the site name, city, and total enrolment count, ordered by count descending.

```
SELECT hs.site_name, hs.city, COUNT(e.enrolment_id) AS total_enrolments
FROM   dbo.HospitalSite hs
LEFT   JOIN dbo.Enrolment e ON e.site_id = hs.site_id
GROUP BY hs.site_name, hs.city
ORDER BY total_enrolments DESC;
```

Screenshot.



	site_name	city	total_enrolments
1	Westbrook General Hospital	London	4
2	St Katherine's Hospital	Manchester	4
3	Ridgeway Community Hospital	Bristol	0

Figure 6: Query 5 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. Starting from `HospitalSite` with a `LEFT JOIN` preserves sites with no enrolments, and `COUNT(e.enrolment_id)` counts rows rather than `NULLs` so empty sites correctly score 0. `Enrolment` stores `site_id` directly, so it is joined straight off `HospitalSite` with no `TrialSite` hop needed.

Query 5bis (companion). The clinical question starts with “how many *participants*” but then specifies a “total *enrolment* count” as the output, and the two differ whenever someone enrolls more than once at the same site. I therefore added a companion query that counts unique individuals instead:

```
SELECT hs.site_name, hs.city, COUNT(DISTINCT e.participant_id) AS
    total_participants
FROM   dbo.HospitalSite hs
LEFT   JOIN dbo.Enrolment e ON e.site_id = hs.site_id
GROUP BY hs.site_name, hs.city
ORDER BY total_participants DESC;
```

	site_name	city	total_participants	
1	Westbrook General Hospital	London	3	
2	St Katherine's Hospital	Manchester	3	
3	Ridgeway Community Hospital	Bristol	0	

Figure 7: Query 5bis output grid (VS Code mssql extension, SQL Server 2025).

Data quality: the two queries return different, equally “correct” numbers from the same data: Query 5 gives Westbrook General and St Katherine’s 4 each, Query 5bis gives 3 each (WBK-002 and WBK-003 each enrolled in a second trial at the same site). This is a classic reporting issue I have faced in my company: teams answering the same loosely worded request publish figures that disagree, and reconciling them afterwards costs time and trust. I always say: *ask twice, report once*, clarify the request before running the query.

Query 6: Trials with an amended protocol

Clinical question. The regulatory affairs team wants to identify trials where the protocol has been amended, meaning more than one consent version exists. Show the trial name, current status, and the number of consent versions issued, ordered by number of versions descending. Only show trials where more than one version has been recorded.

```
SELECT ct.trial_name, ct.status, COUNT(cv.consent_version_id) AS version_count
FROM   dbo.ClinicalTrial ct
JOIN   dbo.ConsentVersion cv ON cv.trial_id = ct.trial_id
GROUP BY ct.trial_name, ct.status
HAVING COUNT(cv.consent_version_id) > 1
ORDER BY version_count DESC;
```

Screenshot.

	trial_name	status	version_count
1	CARDIOPROTECT: Cardiac Rehabilitation in Post...	active	2

Figure 8: Query 6 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. Grouping by trial and counting versions, then I filter only on trials with more than one version. The `GROUP BY` aggregates the rows, and `HAVING COUNT(...) > 1`, isolates amended trials. I used `HAVING` because it filters the aggregate, which `WHERE` cannot.


Data quality: this counts version rows, so a duplicate version entered in error would overstate the number of genuine amendments.

Query 7: Most recently enrolled participant per site

Clinical question. The research governance team wants to identify the most recently enrolled participant at each hospital site across all trials, to prioritise follow-up contact for new joiners. Show the participant study code, trial name, site name, and enrolment date. Only show **the most recently** enrolled participant per site.

```
WITH RankedEnrolments AS (  
    SELECT p.study_code, ct.trial_name, hs.site_name, e.enrolment_date,  
           RANK() OVER (PARTITION BY hs.site_id  
                        ORDER BY e.enrolment_date DESC) AS rn  
    FROM    dbo.Enrolment      e  
    JOIN    dbo.Participant    p ON p.participant_id = e.participant_id  
    JOIN    dbo.ClinicalTrial  ct ON ct.trial_id      = e.trial_id  
    JOIN    dbo.HospitalSite   hs ON hs.site_id      = e.site_id  
)  
SELECT study_code, trial_name, site_name, enrolment_date  
FROM    RankedEnrolments  
WHERE   rn = 1  
ORDER   BY site_name;
```

Screenshot.



	study_code	trial_name	site_name	enrolment_date
1	WBK-005	CARDIOPROTECT: Cardiac Rehabilitation in ...	St Katherine's Hospital	2025-04-01
2	WBK-003	CARDIOPROTECT: Cardiac Rehabilitation in ...	Westbrook General Hospital	2024-10-01

Figure 9: Query 7 output grid (VS Code mssql extension, SQL Server 2025).

Discussion. As seen in class, I used `RANK()` partitioned by site and ordered by enrolment date descending to rank each site's enrolments. Because tied dates share a rank, two participants enrolled at the same site on the same day would both be returned rather than one being picked arbitrarily. `Enrolment` stores `trial_id` and `site_id` directly, so `ClinicalTrial` and `HospitalSite` are joined straight off it with no `TrialSite` hop needed.

Data quality: enrolment is recorded only to the day, so joint “most recent” participants are indistinguishable. A finer timestamp would be needed to single one out such as `DATETIME`.

Part 3: Reflection

Why no participant names?

Storing only a `study_code` pseudonymises the registry, following the data-minimisation principle of GDPR Article 5(1)(c)¹. The Caldicott Principles² request that every use of confidential information is justified and that the minimum necessary should be used. This reinforces Article 5(1)(c). A breach of the trial database would therefore not directly expose identities, and names add no information the registry needs to fulfil its research purpose.

Direct identifiers such as names would live in a separate, access-controlled master linkage list held by the research office. In this design, researchers can query and report entirely on `study_code`, so trial analyses can be shared more freely because they carry no personal identifiers. The cost is potential friction and the dependency on that linkage. A broken mapping could mean a participant cannot be reached. Using a pseudonymised design trades day-to-day convenience for a much stronger confidentiality and information-governance position. This is critical when dealing with personal data.

A situation where the data could mislead

Consent records can mislead because they are true. The scenario requires that when a protocol is amended, every active participant signs the new consent version (*“Each participant must sign the current consent form when they are enrolled. If the protocol changes during their participation and a new consent form is issued, they must sign the new version too.”*).

The database can only validate rows that exist, not rows that are missing. A participant can remain “active” in CARDIOPROTECT with their only signature on version 1, signed before the amendment. Anyone checking whether they consented will find a correctly witnessed record and conclude that they have. Nothing on that row says *superseded*. The registry looks compliant, yet the participant is potentially being treated under wording they never agreed to. The data is not wrong, it is incomplete.

To detect it, I would write a query flagging every active enrolment that has no signing for its trial’s current consent version. I would wrap that query in a view and run it on a schedule with a SQL Server Agent job. I would feed the view into a reporting tool. However, this is **not a database fix**. The database cannot force a missing row into existence. The view and the job only highlight the gap for humans to act on.

Why consent records are append-only, and how to enforce it

Consent records are the legal and ethical evidence that each participant agreed to the specific protocol version in force **at the time**. They must be append-only because consent is a historical fact, and ICH E6 Good Clinical Practice³ requires trial-critical records to carry a complete, tamper-evident audit trail. Editing or deleting a signing event would destroy the evidence that regulators, sponsors and ethics committees rely on to prove valid consent existed. If the rule were broken, a re-consent after a protocol amendment could be silently overwritten or a withdrawn consent erased, making it impossible to demonstrate which wording a participant actually agreed to and undermining every downstream analysis that assumes valid consent.

I enforced this at the database level with an `INSTEAD OF UPDATE, DELETE` trigger on `ConsentSigning` that rejects any such operation with an error, while leaving `INSERT` free (see `script/HDS_DB_LAINSBURY_2026.sql`). This guarantees the rule regardless of which application or user connects, which a permission grant alone would not.

```
CREATE OR ALTER TRIGGER dbo.trg_ConsentSigning_AppendOnly
ON dbo.ConsentSigning
INSTEAD OF UPDATE, DELETE
AS
BEGIN
    SET NOCOUNT ON;
    THROW 51000, 'ConsentSigning is append-only: rows cannot be updated or
    ↪ deleted.', 1;
    -- throw a custom error to reject the operation so in case of an attempt to
    ↪ update or delete, the trigger will fire and allow the operation to be
    ↪ rejected with a clear message.
END;
```

References

1. Art. 5 GDPR – Principles relating to processing of personal data. Regulation (EU) 2016/679. Published online 2016. Accessed July 7, 2026. <https://gdpr-info.eu/art-5-gdpr/>
2. National Data Guardian. The Caldicott Principles. Published online 2020. Accessed July 7, 2026. <https://www.gov.uk/government/publications/the-caldicott-principles>
3. International Council for Harmonisation. ICH E6(R3) Guideline for Good Clinical Practice. Published online 2025. Accessed July 7, 2026. <https://www.ema.europa.eu/en/ich-e6-good-clinical-practice-scientific-guideline>

Appendix A: Data dictionary

A complete column-level reference for every table, derived from the `CREATE TABLE` section of `script/HDS_DB_LAINSBURY_2026.sql`. *Key*: PK = primary key, FK = foreign key, UK = unique (business) key. All surrogate `*_id` keys are `INT IDENTITY(1,1)`.

This dictionary is not static documentation: the metadata section of `script/HDS_DB_LAINSBURY_2026.sql` stores every table description **in the database catalog itself** as `MS_Description` extended properties. The dictionary can therefore be regenerated at any time (or consumed by data-catalogue tooling) by querying the system views, so the metadata travels with the schema rather than living in a document that can drift out of date:

```
SELECT t.name AS table_name,
       c.name AS column_name,
       TYPE_NAME(c.user_type_id) AS data_type,
       IIF(c.is_nullable = 1, 'NULL', 'NOT NULL') AS nullability,
       CAST(ep.value AS NVARCHAR(400)) AS column_description
FROM sys.tables t
JOIN sys.columns c ON c.object_id = t.object_id
LEFT JOIN sys.extended_properties ep
      ON ep.major_id = c.object_id AND ep.minor_id = c.column_id
      AND ep.class = 1 AND ep.name = 'MS_Description'
ORDER BY t.name, c.column_id;
```

ClinicalTrial

Table 5: ClinicalTrial

Column	Type	Null	Key	Notes
trial_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
registration_number	VARCHAR(20)	NOT NULL	UK	ISRCTN number; externally recognised natural key.
trial_name	NVAR- CHAR(150)	NOT NULL		Full trial title.
phase	VARCHAR(4)	NOT NULL		CHECK IN ('I','II','III','IV').
status	VARCHAR(20)	NOT NULL		CHECK IN ('recruiting','active','completed','suspended').

Column	Type	Null	Key	Notes
start_date	DATE	NOT NULL		Trial start date.
end_date	DATE	NULL		NULL = ongoing. CHECK (end_date >= start_date).

HospitalSite

Table 6: HospitalSite

Column	Type	Null	Key	Notes
site_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
site_code	VARCHAR(10)	NOT NULL	UK	Short operational code, e.g. WBK-GEN.
site_name	NVAR- CHAR(100)	NOT NULL		Hospital site name.
city	NVAR- CHAR(60)	NOT NULL		City.
country	NVAR- CHAR(60)	NOT NULL		Country.

Participant

Table 7: Participant

Column	Type	Null	Key	Notes
participant_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
study_code	VAR- CHAR(15)	NOT NULL	UK	Pseudonym; only external identifier (no names held).
date_of_birth	DATE	NOT NULL		CHECK (date_of_birth < registration_date).
sex_at_birth	VAR- CHAR(10)	NOT NULL		CHECK IN ('Male','Female'). Records biological sex at birth, not gender identity — only two values are captured for this field.
registration_date	DATE	NOT NULL		Date first registered with the research office.

InactivationReason

Table 8: InactivationReason

Column	Type	Null	Key	Notes
reason_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
reason_code	VARCHAR(30)	NOT NULL	UK	Standardised report token, e.g. consent_withdrawn.

Column	Type	Null	Key	Notes
reason_description	NVAR-CHAR(255)	NOT NULL		Human-readable description of the reason.

TrialSite (link table)

Table 9: TrialSite. The only table in the schema without a surrogate IDENTITY key — the (trial_id, site_id) pairing is itself the row’s identity.

Column	Type	Null	Key	Notes
trial_id	INT	NOT NULL	PK, FK	-> ClinicalTrial. Part of the composite PK (trial_id, site_id).
site_id	INT	NOT NULL	PK, FK	-> HospitalSite. Part of the composite PK (trial_id, site_id).
opened_date	DATE	NOT NULL		Date the site opened for recruitment on this trial.
closed_date	DATE	NULL		NULL = still open. CHECK (closed_date >= opened_date).

ConsentVersion

Table 10: ConsentVersion

Column	Type	Null	Key	Notes
consent_version_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
trial_id	INT	NOT NULL	FK, UK	-> ClinicalTrial. Part of UNIQUE (trial_id, version_number).
version_number	INT	NOT NULL	UK	CHECK (version_number > 0). Part of the unique key.
effective_from	DATE	NOT NULL		Date this version came into effect.
wording_text	NVAR-CHAR(MAX)	NOT NULL		Full consent-form wording; never overwritten.

Enrolment

Table 11: Enrolment. A filtered unique index allows only one active enrolment per participant (WHERE inactivation_date IS NULL).

Column	Type	Null	Key	Notes
enrolment_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
participant_id	INT	NOT NULL	FK, UK	-> Participant. Part of UNIQUE (participant_id, trial_id, site_id, enrolment_date).

Column	Type	Null	Key	Notes
trial_id	INT	NOT NULL	FK, UK	Composite FK with site_id -> TrialSite (confirms the site actually runs the trial). Part of the unique key.
site_id	INT	NOT NULL	FK, UK	Composite FK with trial_id -> TrialSite . Part of the unique key.
enrolment_date	DATE	NOT NULL	UK	Date enrolment began. Part of the unique key (same-day duplicate-entry guard).
inactivation_date	DATE	NULL		NULL = still active. CHECK (inactivation_date >= enrolment_date).
inactivation_reason_id	INT	NULL	FK	-> InactivationReason . Paired-null CHECK with inactivation_date .

ConsentSigning (append-only)

Table 12: ConsentSigning. Protected by an **INSTEAD OF UPDATE, DELETE** trigger so rows can only ever be inserted.

Column	Type	Null	Key	Notes
consent_signing_id	INT IDENTITY	NOT NULL	PK	Surrogate primary key.
enrolment_id	INT	NOT NULL	FK, UK	-> Enrolment . Part of UNIQUE (enrolment_id , consent_version_id , signed_date).
consent_version_id	INT	NOT NULL	FK, UK	-> ConsentVersion . Part of the unique key.
signed_date	DATE	NOT NULL	UK	Date the consent was signed. Part of the unique key (same-day duplicate-entry guard).
witnessed_by	NVAR- CHAR(100)	NOT NULL		Name/role of the witnessing staff member.